

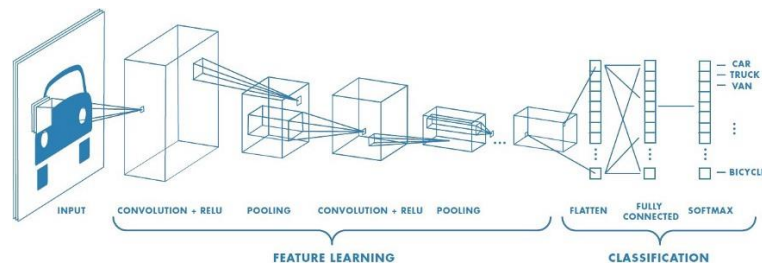
Convolutional Neural Network (CNN)

Reference:

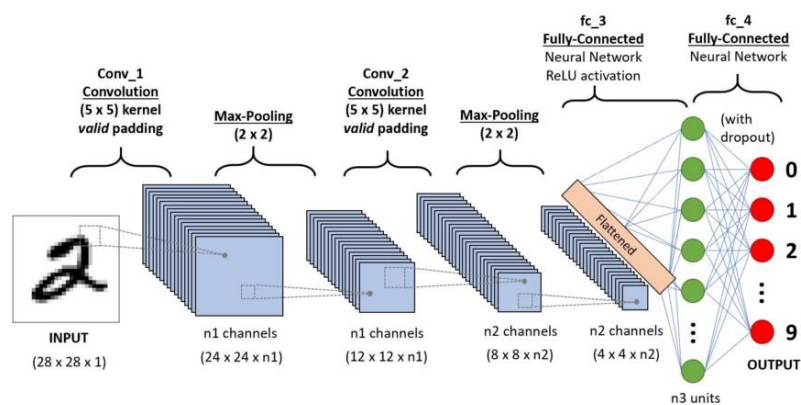
A comprehensive guide to CNN, <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

Yangfeng Ji, Machine learning, lecture notes

1 Overview



- CNN works well for image classification (benefits from **feature learning** algorithms)
- This architecture repeats the two components twice to learn features in images before connecting with a fully-connected classification task.
 - 1) convolutional layer
 - 2) subsampling (pooling) layer
- The main steps of CNN
 - 1) Rescale data into suitable range $[0,1]$
 - 2) Feature learning (twice): convolution+ReLU+pool
 - 3) Flat the features into 1-dimension after step 2
 - 4) Fully-connected Neural Network
 - 5) **Notice**: pay attention to the dimension of data in the process




2 Why ConvNets over Feed-Forward Neural Nets?

- **Feed-Forward Neural Nets:** flat the image and feed it into multi-level perception (MLP). However, it will lose relative information.
- **ConvNet:** For complex images with pixel dependencies throughout, it can successfully **capture the Spatial and Temporal dependencies** in an image through the application of relevant filters (Convolution+ReLU+Pooling). The architecture performs a better fitting to the image dataset by **reusing weights and reducing the number of parameters**.

3 Input Image

- **Images:** can be grayscale (1 channel), RGB (3 channel). Furthermore, the images can be complex.
- **Function of ConvNet:** reduce the images into a form that is easier to process, without losing features that are critical for getting a good prediction.

4 Convolution Layer — The Kernel




- The **input image**: has dimension $n_1 \times m \times m$, where n_1 is the number of channels, $m \times m$ is the size of the image for a single channel.
- **The kernel**: should also have the same number of channels, i.e., $n_1 \times k \times k$.
- **The output**: the dimension is affected by the input image, kernel, and Stride_length.

In the above demonstration, the green section resembles our 5x5 input image, I. The element involved in carrying out the convolution operation in the first part of a Convolutional Layer is called the **Kernel/Filter, K**, represented in the color yellow. We have selected **K as a 3x3 matrix**.

```
Kernel/Filter, K =
1 0 1
0 1 0
1 0 1
```

The Kernel shifts 9 times because of Stride Length = 1 (Non-Strided), every time performing a matrix multiplication operation between K and the portion P of the image over which the kernel is hovering.

- **Kernel function in Pytorch:** (the following kernel with padding=1)



CLASS `torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None)` [SOURCE]


Parameters:

- **in_channels** (*int*) – Number of channels in the input image
- **out_channels** (*int*) – Number of channels produced by the convolution
- **kernel_size** (*int or tuple*) – Size of the convolving kernel
- **stride** (*int or tuple, optional*) – Stride of the convolution. Default: 1
- **padding** (*int, tuple or str, optional*) – Padding added to all four sides of the input. Default: 0
- **padding_mode** (*str, optional*) – 'zeros', 'reflect', 'replicate' or 'circular'. Default: 'zeros'
- **dilation** (*int or tuple, optional*) – Spacing between kernel elements. Default: 1
- **groups** (*int, optional*) – Number of blocked connections from input channels to output channels. Default: 1
- **bias** (*bool, optional*) – If `True`, adds a learnable bias to the output. Default: `True`

<https://pytorch.org/docs/stable/generated/torch.nn.Conv2d.html>

5 Pooling Layer

- **Functions of Pooling layer:** i) Reduce the spatial size of the Convolved Feature. ii) extracting dominant features.
- **Pool types:** 1) Average pooling; 2) Max pooling. In particular, Max Pooling also performs as a Noise Suppressant. we can say that **Max Pooling performs a lot better than Average Pooling**.

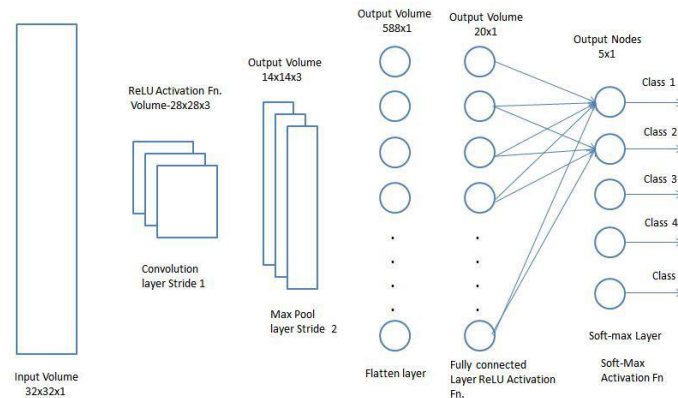


6 A Complete ConvNet layer

The **Convolutional Layer** and the **Pooling Layer**, together form the i -th layer of a Convolutional Neural Network.

After all complete **convolution Nets**, we are going to **flatten the final output** and feed it to a regular Neural Network for classification purposes.

7 Classification — Fully Connected Layer (FC Layer)



Adding a Fully-Connected layer is a (usually) cheap way of learning **non-linear combinations** of the high-level features as represented by the output of the convolutional layer.

Then the fully-connected MLP is **trained** by feeding **the flattened output** to a feed-forward neural network and **backpropagation** is applied to every iteration of training.

Over a series of epochs, the model is able to distinguish between dominating and certain low-level features in images and classify them using the **Softmax Classification** technique.

8 Example-CNN for the Fashion Minist dataset

- Input data: $1 \times 28 \times 28$. The batch size $S = 20$
- The CNN network

```
class Net(nn.Module): # Set convolutional network
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(20, 32, 5)
        self.fc1 = nn.Linear(32 * 4 * 4, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

input channel output channel
 $\text{Conv2d}(d, b, c)$ kernel size
 $1 \times 28 \times 28 \rightarrow (1, 20, 5 \times 5)$
 $\rightarrow 20 \times 24 \times 24 \xrightarrow{\text{max pool}(2,2)} 20 \times 12 \times 12$
 $\rightarrow 20 \times 12 \times 12 \xrightarrow{\text{kernel}(20, 32, 5 \times 5)} 32 \times 8 \times 8$
 $\rightarrow 32 \times 8 \times 8 \xrightarrow{\text{max pool}(2,2)} 32 \times 4 \times 4$
 $\rightarrow \text{flatten}$

- **Cost function and optimization method:** CrossEntropyLoss and SGD
<https://pytorch.org/docs/stable/optim.html>

The `torch.optim` function: To use `torch.optim` you have to construct an optimizer object, that will hold the current state and will update the parameters based on the computed gradients.

```
import torch.optim as optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)

# #####
```

```
# zero the parameter gradients
optimizer.zero_grad()

# forward + backward + optimize
outputs = net(inputs)
loss = criterion(outputs, labels)
loss.backward()
optimizer.step()
```

- The main **optimization steps** of one round includes:
 - a) Clean the gradient,
 - b) generate the output of a model by a given input,
 - c) Compute the loss
 - d) Backpropagation the gradient
 - e) update the model by one optimization step

- The classification accuracy of CNN: 90%

```
GroundTruth: Ankle boot Pullover Trouser Trouser
Predicted: Ankle boot Pullover Trouser Trouser
Accuracy of the network on the 10000 test images: 90 %
Accuracy for class: T-shirt/top is 85.3 %
Accuracy for class: Trouser is 97.9 %
Accuracy for class: Pullover is 88.6 %
Accuracy for class: Dress is 90.9 %
Accuracy for class: Coat is 90.6 %
Accuracy for class: Sandal is 97.7 %
Accuracy for class: Shirt is 68.2 %
Accuracy for class: Sneaker is 97.7 %
Accuracy for class: Bag is 98.3 %
Accuracy for class: Ankle boot is 94.2 %
```